

# **WEIGHT IN THE ATTENTION**



**ISSUE 202508**

**PREPRINT**

**The founding of Weight in the Attention,  
a journal for the AIGC Zeitgeist P.1**


**WITA Manuscript Submission Guide P.3**



# The founding of Weight in the Attention, a journal for the AIGC Zeitgeist

Neruthes (WITA Editor)

2025-07-17

 Article copyright licensing scheme: **CC BY-ND 4.0**.

## Abstract

This article is a founding memo of the Weight in the Attention journal, a project inspired by the recent AIGC bubble. In time of post-COVID economic stringency, the AIGC landscape presents a dramatic spectacle of eager to recreate the good old pre-COVID entrepreneurship hype. Existing works [3] have demonstrated a promising vision for making fun of the time and the trend as both an outsider and an insider. The author examined possible vacancies for AIGC-oriented humor and satire and proposed the creation of this journal as a leisure platform.

## 1 Zeal, FOMO, and Anxiety

The introduction of DeepSeek in early 2025 put an extra burden [5] on already stressed LLM engineers. Difficulties were observed on enjoying shared moments with friends who were working at its competitors. In a highly competitive environment with little regulatory determination on maintaining work-life balance [2] [4], the fear of missing out (FOMO) could easily pass down from top executives all the way along to frontline product managers and software engineers, regardless of working on training models, tweaking contexts, or developing supportive systems.

## 2 Buzzwords, Burnout, and Depression

On the other hand, being out of the AIGC industry does not mean feeling better. Since the release of ChatGPT, buzzwords proliferated one-after-another — transformer, decoder,

attention, agents, context, RAG, etc. The landscape looked like the radical evolution of web frontend during 2014-2020 after with radical abandon of jQuery — Backbone, Meteor, Vue, Angular, React, Gulp, Grunt, WebPack, Next.js, etc. The quantity of technical stack choices in web frontend projects were doubling every 18 months [1].

Both trends may lead to the same consequence — burnout. I had personally experienced the burnout with the radical evolution of web frontend, and later decided to stick to good old vanilla approach for small projects of personal use and to wait till a promising silver bullet emerges to give the zeal a pause, otherwise I would have no cognitive resource to keep learning new web frontend stuff especially when it was not my main focus. No person is absolutely immune to burnout and attention is not infinite. To keep up to date does not constitute to follow every tiny footstep without examining its prowess to remain part of best practices in subsequent years given that one does not personally earn a living from it. In hindsight, the moments suggested clues of depression.

### 3 Content Orientation

This journal shall welcome articles that fall under any of the following categories.

1. **Written by LLM:** The article is mainly written by LLM.
2. **Social impact:** The article studies social impacts of AIGC such as financial bubble, geopolitical contest, and end-user reception.
3. **Interdisciplinary study:** The article discusses how AIGC interacts with other studies.
4. **Usage best practices:** The article affords information on how to develop, deploy, or prompt AIGC software products.
5. **Onion news:** The article humorously resembles the absurdities of this era with fake facts and true insights.

### References

- [1] G.E. Moore. “Cramming More Components Onto Integrated Circuits”. In: *Proceedings of the IEEE* 86.1 (Jan. 1998), pp. 82–85. ISSN: 1558-2256. DOI: 10.1109/JPR0C.1998.658762.
- [2] X. Zheng and Z. Qiu. “The 996 working pattern in Chinese internet firms: How hegemonic despotism promotes long working hours for employees”. In: *China Perspectives* 134.1 (2023), pp. 67–78. ISSN: 1996-4617. DOI: 10.4000/chinaperspectives.15869.
- [3] *AGI Bar* 知识蒸馏. 2025. URL: <https://agi.bar/>.
- [4] Ming Liu and Yunqiao Chen. “Blessing or curse? Recontextualizing ‘996’ in China’s overwork debate”. In: *Critical Discourse Studies* 22.1 (2025), pp. 91–107. DOI: 10.1080/17405904.2023.2289448. eprint: <https://doi.org/10.1080/17405904.2023.2289448>. URL: <https://doi.org/10.1080/17405904.2023.2289448>.

- [5] Tony Peng. *RedNote Investor on DeepSeek: The Next Android of the AI Era*. 2025. URL: <https://recodechinaai.substack.com/p/rednote-investor-on-deepseek-the>.

# WITA Manuscript Submission Guide

Neruthes (WITA Editor)

2025-07-14

**i** Article copyright licensing scheme: **CC BY-ND 4.0**.

## Abstract

Like most journals, Weight in the Attention (WITA) handles manuscript submissions. A rare characteristic of WITA is that the submission process is based on Git, GitHub, and pull request. This guide offers a comprehensive guide for authors in good faith of establishing an efficient manuscript acceptance workflow. Key takeaway — if you would like to submit an already published article, just open an issue and include URL to your article.

## 1 Directory Structure

Issues are grouped by year inside the “/issue” directory. Each issue consists of a “tex” file and a “tex.d” directory.

### 1.1 Issues Catalog

Sketch of directory structure “/issue”:

```
/issue/  
|- 2025/                -> Year, format YYYY  
|   |- 202508.tex       -> IssueID, format YYYYMM  
|   |- 202508.tex.d/  
|       |- meta/  
|       |- entry/
```

## 1.2 Inside Single Issue

Sketch of directory structure “/issue/Year/IssueID.tex.d”:

```
../202508.tex.d/      -> IssueDir
| - meta/
| - entry/
|   - 001/            -> EntryID
```

## 2 Organizing Your Entry

Sketch of directory structure “/issue/Year/IssueID.tex.d/entry/EntryID”:

```
../001/              -> EntryDir
| - info.toml
| - main.tex
| - cite.bib
```

You should have these 3 files. In addition, you may have subdirectories for images and code pieces.

### 2.1 info.toml

This is a TOML file that contains the metadata of the submitted manuscript.

Example:

```
[[article]]
id = "myrun"                # Same to EntryID
title = "Article Title Goes Here"
authors = ["John Appleseed (Reed College)", "John Doe (*)"]
authors_simple = "Appleseed, J., et al."
email = "user@example.com"
date = "2025-07-13"
license = "CC BY-ND 4.0"    # Omit NC to allow more
```

### 2.2 main.tex

This is where your article content goes.

Example:

```
\setentryid{myrun}
\stdarticle{Article Title Goes Here}{%
  \authorrow{John Appleseed}{(Reed College)}\\%
  \authorrow{John Doe}{}%
}{2025-07-13}
```

You should start the file like the example.

Here are commands you can use:

- `\setentryid`  
Declares article EntryID.
- `\stdarticle{Title}{Authors}{SubmissionDate}`  
Print article title, author, and date information.
- `\entrypath`  
Get path to entry. Useful for `\includegraphics{\entrypath/pic-1}`.
- `\authorrow{Name}{Institution}`  
Prints a row of author information. Must use inside `argv1` of `\stdarticle`.

## 2.3 cite.bib

Unlike others, **cite.bib** may be omitted if your manuscript contains no citation at all.

## 3 Submission Workflow

Follow these steps to submit your article to WITA.

1. **Fork Repository:** Fork the repository and clone your fork to your machine.
2. **Elect EntryID:** Decide a short string for your article that is unlikely to collide with other authors of the same entry.
3. **Create Branch:** In your local repository, create a branch named after your username (e.g. “johndoe/entry256”).
4. **Create Entry:** In your branch, create relevant directories and files.
5. **Local Build:** Run the issue building workflow locally and debug problems, if any.
6. **Create Pull Request:** Create a pull request that merges from your fork “johndoe:johndoe/entry256” to upstream master branch.
7. **Await Editor Review:** Editors will review your article and will very likely accept it.

The following list contains further clarification.



1. You are free to, and encouraged to, self-publish your article anywhere else such as your personal blog. It is *your* article, after all.
2. If you include non-empty “email” field in your “info.toml” file, an editor will mail you a letter of acceptance when your article is accepted.
3. Even if your article is accepted, there is no guarantee that it will be included in the next coming issue.
4. When creating pull request, use dummy IssueID “000000” (/issue/0000/000000.tex.d). An editor will move your article to an upcoming IssueID after acceptance.
5. To check whether your article is bug-free, you can run `./make.sh issue/0000/000000.tex`. If you do not have a GNU/Linux machine, try WSL.
6. Your article should have a redistributable license such, e.g. CC BY-ND, CC BY-SA, GFDL. Alternatively, you can make it public domain. Open knowledge is great.
7. LLM-created articles are encouraged to be submitted as public domain work.
8. If the article is created by LLM, the LLM should be attributed as an author. Name can be online public service (e.g. Gemini, ChatGPT) or model identifiers (e.g. Qwen3-32B).

# My Startup Is Just a Claude Wrapper, but It's a Free Claude Wrapper

Neruthes  
ChatGPT (OpenAI)

2025-07-14

📄 Article copyright licensing scheme: **Public Domain**.

## Abstract

In 2025's AI frenzy, few brag sheets boast more bravado than “We're a *Claude wrapper*.” At conferences, hackathons, and booster-pitch dinners, someone inevitably leads with: “We didn't train a model — we integrated Claude's API in 30 minutes, wrapped a lightweight UI, and voilà!” Congratulations: your startup is officially “innovative”—as long as flimsy UI counts as disruptive.

## 1 The Rise of the Wrapper

A decade ago, “wrapper” meant a bandage; now it's a business model. A popular observation calls these *lightweight applications* built on third-party LLM APIs, with minimal effort and complexity [7]. Indeed, startups like Manus AI—which bridges Claude with dozens of tools—are marketed as transformative, yet critics note: “Technically, yes — Manus uses Claude's API connected to 29 different tools” [3].

Yet that hasn't stopped the cash flow: Bolt.new, a Claude-based code agent, reportedly pulled in *\$8 million ARR in two months* [9]. On Reddit and LinkedIn, the common refrain is: “they may be wrappers, but their UI/bundling is what sells” [8]. In short, convenience *is* value — if users pay.

## 2 Why Wrappers Sell — Even If They’re Middling

### 2.1 Speed to Market

You don’t need to train a model on petabytes of data — just spin up OpenAI’s or Anthropic’s API, slap on a dashboard, and you’re live.

### 2.2 Product-Market Fit

Fast iteration beats slow perfection. A wrapper with slightly better UX often trumps slower in-house models [5, 8].

### 2.3 Hype-Fueled Funding

Startup valuations are still riding the AI bubble. Investors throw money at any venture tagged *AI*, even if it’s just an API client.

Is this sustainable? Some argue that *99% of AI startups will die by 2026*—not due to fraud, but because wrappers lack defensible advantages, product moats, or infrastructure [5].

## 3 The Good, the Bland, and the Ugly

### 3.1 The Good

Wrappers democratize access. Tools like Manus or Lumio AI let non-coders assemble workflows, compare models, and embed intelligence into documents or code — no PhD required [6].

### 3.2 The Bland

Yet a face-emoji web app or generic chatbot? Meh. Many wrappers add no real power—just paint. They echo the “AI Snake Oil” critiques by Narayanan, who warns that marketing often overshadows technical merit [1].

### 3.3 The Ugly

Bandwagon effects hide deeper issues: pumps & dumps of attention, lack of security, poor UX, bias inherited via opaque APIs, and zero explainability. The prompt-engineer-as-founder has become a punchline—and a weak business model [2].

## 4 What We Lose If Wrappers Win

- **Lack of innovation:** Relying on wrappers stifles model-level breakthroughs.
- **Commoditization:** AI becomes a plug-and-play feature, not a transformative platform.
- **Dependency risks:** Lock-in to API ecosystems, vendor price hikes, and downstream fragility.
- **Ethical opacity:** Wrappers perpetuate bias and hallucination without oversight [10].

## 5 Where to Go from Here

- *Layer up, don't wrap:* Build meaningful differentiation—team curation, proprietary finetunes, UX innovations, deep data integration.
- *Prove real ROI:* Move past MVPs to revenue, retention, or B2B traction—beyond Buzzword Bingo.
- *Push transparency:* Embrace explainability frameworks, model audits, and bias remediation.
- *Prepare for shakeout:* Many wrapper ventures will fade; the survivors will innovate, not imitate [4].

## 6 Final Word: Wrappers Can Blossom or Burst

Calling yourself a “Claude wrapper” isn’t shameful — but it shouldn’t stop at calling. Execution, defensible strategy, and genuine impact matter. The current hype bubble may lift all boats — but once the tide recedes, only the ships that actually sail will remain.

So yes — your startup might technically be a “free Claude wrapper.” But is it just dressing, or will it deliver substance? That’s the question investors, users, and founders must answer — without the marketing smoke.

## References


- [1] Arvind Narayanan and Sayash Kapoor. “AI Snake Oil”. In: *AI Snake Oil Blog* (2023). URL: <https://aisnakeoil.substack.com>.
- [2] Dion Wiggins. *The Real-Life AI Hype Curve*. LinkedIn. 2024. URL: <https://www.linkedin.com/pulse/real-life-ai-hype-cycle-aka-burn-curve-dion-wiggins-wx4kc>.

- [3] Julian Goldie. *Why Manus AI Is Making Me Rethink My Entire Business Model*. Medium. 2025. URL: <https://medium.com/@julian.goldie/why-manus-ai-is-making-me-rethink-my-entire-business-model-and-you-should-too-c3a21be61fc6>.
- [4] *Is Anyone Actually Making Something Which Is Not Just a Wrapper?* Reddit /r/SaaS. 2025. URL: [https://www.reddit.com/r/SaaS/comments/1hdbt3q/is\\_anyone\\_actually\\_making\\_something\\_which\\_is\\_not/](https://www.reddit.com/r/SaaS/comments/1hdbt3q/is_anyone_actually_making_something_which_is_not/).
- [5] Skool of Life. *99% of AI Startups Will Be Dead by 2026—Here's Why*. Medium. 2025. URL: <https://skooloflife.medium.com/99-of-ai-startups-will-be-dead-by-2026-heres-why-bfc974edd968>.
- [6] *Lumio AI*. Wikipedia entry. 2025. URL: [https://en.wikipedia.org/wiki/Lumio\\_AI](https://en.wikipedia.org/wiki/Lumio_AI).
- [7] Zack Olivas. “*AI Startups Are Just Wrappers*”. LinkedIn Post. 2025. URL: [https://www.linkedin.com/posts/zack-olivas\\_i-keep-seeing-that-ai-startups-are-all-chatgpt-activity-7254229614699450369-QuZm](https://www.linkedin.com/posts/zack-olivas_i-keep-seeing-that-ai-startups-are-all-chatgpt-activity-7254229614699450369-QuZm).
- [8] Krunal Patel. *Unwrapping the Hype: Drawbacks of Wrapper Startups*. Blog. 2025. URL: <https://krunal.org/unwrapping-the-hype-the-drawbacks-of-wrapper-startups-in-ai-170f973040c6>.
- [9] Latent Space. *Bolt: \$8M ARR Claude Wrapper*. Podcast and Blog. 2025. URL: <https://www.latent.space/p/bolt>.
- [10] *The AI Revolution Is Already Losing Steam*. Wall Street Journal. 2025. URL: <https://www.wsj.com/tech/ai/the-ai-revolution-is-already-losing-steam-a93478b1>.

# Navigating the AI Geoscape: A Multi-Faceted Analysis of Manus AI's Strategic Retreat from Mainland China

Neruthes  
Gemini (Google)

2025-07-14

 Article copyright licensing scheme: **Public Domain.**

## Abstract

This report analyzes the recent strategic pivot of Manus AI, a prominent Chinese AI agent startup, involving the cessation of its mainland China operations and the relocation of its global headquarters to Singapore. While geopolitical tensions, particularly US investment restrictions and export controls on advanced AI chips, are widely cited as primary drivers, this analysis argues that Manus AI's move was also significantly influenced by intense domestic market competition, challenges in product differentiation, and a proactive global talent strategy. By examining the interplay of these external and internal pressures, the report offers a holistic perspective on the complex decision-making processes of Chinese AI firms navigating an increasingly polarized global technology landscape, highlighting Singapore's role as a strategic "third path" hub.

## 1 Introduction

Manus AI, developed by Beijing Butterfly Effect Technology, emerged onto the global artificial intelligence (AI) scene in March 2025 with an invite-only AI agent that quickly garnered attention [10]. Its core proposition marked a significant departure from conventional conversational AI, as it was designed to autonomously perform complex, multi-step tasks such as filtering resumes, analyzing stocks, and even writing and deploying code [6]. This capability positioned Manus AI as a general-purpose agent aiming to bridge the gap between human intention and tangible action, functioning as a self-directed digital assistant. The platform generated considerable market buzz, with activa-

tion codes reportedly selling for as much as 100,000 yuan on secondary markets, underscoring the initial excitement surrounding its capabilities.

However, by July 2025, Manus AI undertook a significant, largely silent, operational adjustment, ceasing most of its activities in mainland China and relocating its global headquarters to Singapore [4]. This strategic shift involved substantial workforce reductions in China, with the majority of its approximately 120 staff laid off, while a core group of around 40 technical personnel were transferred to the new Singapore headquarters [7]. Concurrently, Manus AI’s digital footprint in China was erased, with official Chinese social media accounts (Weibo, Xiaohongshu) cleared of content and its official website displaying a message stating “Manus is not available in your region” for Chinese users [2]. This marked a stark change from its previous message indicating a “Chinese version is under development”.

This report posits that Manus AI’s strategic pivot was not merely a reaction to escalating US-China geopolitical tensions but a calculated, multi-faceted response driven by a complex interplay of external pressures, including US investment restrictions and chip export controls, and significant internal market dynamics, such as hyper-competition, declining user engagement, and a proactive global talent strategy.

## **2 The Strategic Pivot: Confirmation and Context**

The operational adjustments undertaken by Manus AI in mid-2025 demonstrate a decisive move away from its original base in mainland China. The company officially moved its headquarters from China to Singapore in June 2025 [5, 12, 7, 8, 11]. This relocation was publicly confirmed by co-founder and chief product officer Zhang Tao at the SuperAI conference in Singapore on June 18, 2025, where he stated that Singapore was now Manus AI’s main base. Beyond Singapore, the company also established offices in Tokyo and California, specifically San Mateo.

A significant component of this pivot was the restructuring of its workforce. Beijing Butterfly Effect Technology, the Chinese operating entity behind Manus AI, laid off most of its approximately 120 employees in China, with those remaining receiving severance packages. Crucially, approximately 40 core technical personnel were transferred to the new Singapore headquarters, indicating a strategic effort to retain essential expertise while shifting the operational center of gravity.

The company’s digital presence also underwent a dramatic transformation. Manus AI’s official accounts on prominent Chinese social media platforms, Weibo and Xiaohongshu, were cleared of all content. Furthermore, a previously announced partnership with Alibaba Cloud’s chatbot, Tongyi Qianwen, was deleted, and a former employee confirmed the collaboration would not proceed. Perhaps most indicative of the shift, the official Manus AI website, which once promised a “Chinese version is under development,” now displays “Manus is not available in your region” for users attempting to access it from China.

In its official communications, Manus AI stated that these adjustments were “based on the company’s own operating efficiency considerations” and aimed to “continue to focus on core business development and improve overall operational efficiency” [14].

The comprehensive nature of these changes, particularly the clearing of Chinese social media accounts and the explicit message blocking access for Chinese users, points to a deliberate strategy to reshape the brand’s identity. This goes beyond mere operational restructuring; it suggests a conscious effort to re-brand Manus AI as a global, non-Chinese entity. The objective appears to be to appeal to international investors and markets, especially in the US, where affiliations with China are increasingly perceived as a liability. This strategic decision to shed its Chinese identity is a clear attempt to mitigate perceived geopolitical risks and enhance its standing in Western markets.

Moreover, the company’s official statement about “operational efficiency” serves as a broad, corporate-friendly explanation for a complex and abrupt withdrawal. While efficiency is a legitimate business objective, the scale and suddenness of Manus AI’s pivot, combined with the explicit geopolitical and market pressures detailed in other reports, suggest that “operational efficiency” functions as a euphemism. This allows the company to depoliticize its actions and avoid publicly acknowledging the full extent of external regulatory pressures and internal market struggles, thereby maintaining a neutral business narrative in a sensitive environment.

Table 1: Key Events and Dates in Manus AI’s China Operations Adjustment

Date	Event Description
March 2025	Manus AI global debut/launch
March	Manus AI monthly active users peak at 20 million
April	Benchmark’s US\$75 million funding round, valuing Manus at US\$500 million
May	Founders relocate to Singapore
May	Manus AI monthly active users fall to 10 million
May	US Treasury Department reportedly reviewing Benchmark funding
June 18	Co-founder Zhang Tao confirms Singapore as HQ at SuperAI conference
Mid-June	Manus ads begin appearing in Singapore
July 9-11	News reports confirm HQ shift, China job cuts, and social media changes

### 3 Geopolitical Pressures: The Dominant Narrative

The most frequently cited catalysts for Manus AI’s relocation are the escalating geopolitical tensions between the United States and China, particularly concerning advanced technology. These external pressures have created a challenging operating environment for Chinese AI firms seeking global reach.



The Biden Administration’s outbound investment regulations, implemented in January 2025, represent a significant constraint on US investors interested in advanced Chinese technology companies. These rules specifically target US investments in Chinese entities involved in sensitive technologies, including AI, by either prohibiting certain transactions or requiring notification to the US Treasury Department. This legislative framework imposes substantial compliance risks and heightened due diligence requirements for US investors, including private equity and venture capital funds. Manus AI’s US\$75 million funding round in April 2025, led by the Silicon Valley venture capital firm Benchmark, immediately drew scrutiny from the US Treasury Department, which initiated a review to determine its compliance with these new restrictions. This demonstrates how investment capital itself has become a tool in geopolitical competition, with the US government actively seeking to prevent its capital from strengthening Chinese AI capabilities perceived as national security risks. The need for external capital for rapid growth in the AI sector is critical, and by relocating, Manus AI attempts to de-risk future fundraising efforts from American investors, illustrating how geopolitical considerations directly influence funding structures and corporate geography.

Complementing investment restrictions are the stringent US export controls on advanced AI chips. Imposed in April 2025, these controls ban the sale of high-end AI chips, such as Nvidia’s H100, to China [8]. Earlier restrictions in October 2022 and 2023 had already targeted Nvidia’s A100 and H100, and even customized lower-performance variants like the H800 and A800. These advanced chips are indispensable for training the complex algorithms that power general-purpose AI agents like Manus AI. Consequently, Chinese firms face considerable obstacles in acquiring these essential components. Nvidia’s CEO, Jensen Huang, has publicly criticized these controls, labeling them a “failure” for inadvertently spurring Chinese companies to accelerate their own AI development and noting a significant decline in Nvidia’s market share in China [15, 3].

In this challenging environment, Singapore has rapidly emerged as a strategically important hub for Chinese-origin technology companies seeking to navigate the tensions between Washington and Beijing. The relocation to Singapore is explicitly aimed at mitigating the impact of US investment restrictions and the escalating US-China AI competition. Singapore offers a compelling value proposition: better access to international markets, crucial computing resources, and global capital. It also provides a pathway for companies to attract Western clients and investors while sidestepping potential restrictions that might otherwise be imposed on China-based entities. This strategic maneuver is not unique to Manus AI; other Chinese tech giants, such as the fast fashion company Shein and the social media platform TikTok, have similarly emphasized their Singapore headquarters while maintaining production networks or control links in China. This trend exemplifies a broader strategy for Chinese AI companies to establish a “third path” in the increasingly polarized US-China tech ecosystem. Singapore, in this context, offers a neutral, internationally connected environment that allows these companies to maintain proximity to China’s talent pool while presenting themselves as global entities, thereby bypassing direct US scrutiny and accessing critical inputs like advanced chips and venture capital. This represents a strategic arbitrage of geopolitical friction.

## **4 Beyond Geopolitics: Internal Market Dynamics and Strategic Repositioning**

While geopolitical pressures undeniably exerted significant influence, Manus AI's strategic pivot was also profoundly shaped by internal market dynamics and a proactive repositioning strategy.

### **4.1 The “War of a Hundred Models”: Intense Domestic Competition**

China's domestic AI market is characterized by an extraordinary level of competition, often referred to as the “war of a hundred models.” With over 130 large language models (LLMs) developed, China accounts for approximately 40% of the global total. This proliferation has led to an intensely crowded market, raising concerns about sustainability and the viability of numerous players. The fierce competition has ignited a “price war,” with major Chinese tech giants like ByteDance, Alibaba, and Baidu drastically cutting prices on their LLM-based services in a bid to attract users. This aggressive pricing environment makes it exceedingly difficult for smaller startups, like Manus AI, to establish viable business models and achieve profitability within the domestic arena. The competitive landscape is further intensified by the entry of established tech giants, such as ByteDance with its Coze Space and Baidu with its AgentBuilder, which have introduced rival AI products, directly competing for market share within China's burgeoning AI sector. This “war of a hundred models” is not merely a challenge; it acts as a powerful internal force pushing Chinese AI startups outwards. While geopolitical factors create external barriers, the intense domestic competition, characterized by a proliferation of similar models and aggressive price wars, makes the Chinese market less attractive for sustainable growth, especially for smaller players. This internal pressure to seek more favorable profit potential internationally complements and amplifies the external geopolitical push, making internationalization a dual imperative for survival and expansion.

### **4.2 Product Viability and User Engagement Challenges**

Manus AI also faced significant challenges related to its product's viability and user engagement. The company experienced a notable decline in its user base, with monthly active users plummeting from approximately 20 million in March to around 10 million by May 2025. This decline suggests underlying issues beyond external pressures.

Despite initial hype, critics have argued that Manus AI lacks “real technological breakthroughs,” describing it as more of a “shell” that relies heavily on existing large models, such as Anthropic's Claude family and Alibaba's Qwen models, and pre-existing toolchains, rather than developing original core technology. This reliance on third-party LLMs contributes to higher operational costs and raises questions about scalability. Fur-

thermore, early feedback on Manus AI highlighted technical issues, including reports of system instability, frequent crashes, inaccurate data generation, and slower processing speeds compared to some competitors. Users also noted difficulties with straightforward operations and integration issues within its multi-agent system. The critique that Manus AI acts as a “shell” relying on existing LLMs points to a deeper challenge in China’s AI ecosystem beyond just quantity. In a market saturated with LLMs, true differentiation requires novel technological advancements. If Manus AI is perceived as lacking originality and merely integrating third-party models, its declining user numbers become understandable, as users may gravitate towards offerings from larger players with deeper research and development capabilities or more unique features. This suggests that the domestic market’s “war of a hundred models” is also a “war of innovation,” where only truly differentiated or highly efficient applications can thrive. The relocation might also be an attempt to access a global talent pool and research and development environment more conducive to developing proprietary, breakthrough technology [16].

### **4.3 Global Talent Strategy and Operational Efficiency**

Manus AI’s pivot also reflects a deliberate global talent strategy and a drive for operational efficiency. The company has aggressively begun recruiting new talent in Singapore, with job listings for positions such as data analyst and AI agent engineer, and is also hiring in the US and Japan. This robust recruitment drive signifies an intention to build a strong international talent base.

Crucially, the strategic transfer of around 40 core technical personnel from China to Singapore ensures the retention of critical research and development capabilities while physically relocating them to a more favorable operating environment. In the highly competitive AI talent market, especially for specialized expertise, retaining key engineers and data scientists is paramount. By relocating them to Singapore, Manus AI can leverage Chinese engineering talent while positioning itself in a global hub that offers better access to international markets, computing resources, and potentially a less restrictive research and development environment. This allows the company to maintain its technical backbone while shedding the geopolitical baggage associated with being fully China-based. Manus AI’s official statements about enhancing “operational efficiency” align with the need to streamline operations in a less competitive, more globally accessible environment. The founder, Zhang Tao, also indicated considering a “full separation of its China and international operations”, suggesting a long-term strategy to completely de-link the global Manus AI brand from its Chinese origins.

## **5 Interplay of Factors: A Holistic Perspective**

Manus AI’s strategic retreat from mainland China was not a singular response to one dominant factor but rather a complex, multi-pronged maneuver necessitated by the synergistic pressures of geopolitical constraints and internal market dynamics. The external

pressures, particularly the US outbound investment regulations implemented in January 2025, significantly limited Manus AI's access to crucial Western capital. This made it increasingly challenging for the company to sustain its operations and compete effectively in the capital-intensive "war of a hundred models" within China. Simultaneously, the US chip export controls, particularly those imposed in April 2025, directly impacted Manus AI's ability to acquire the advanced computing resources essential for training its sophisticated AI algorithms. This fundamental requirement for an AI agent startup was severely hindered, further impeding its technological competitiveness in an already crowded domestic market. The confluence of these external pressures with the internal market saturation and the company's declining user numbers created an untenable operating environment in mainland China, rendering the strategic pivot a matter of both survival and future growth.

Manus AI's actions illustrate the "geopolitical tax" that Chinese tech companies face, forcing them to incur significant costs, including layoffs, relocation, and brand re-orientation, simply to operate globally. The comprehensive nature of Manus AI's withdrawal—encompassing headquarters relocation, mass layoffs, and the erasure of its digital footprint—and its explicit aim to reduce "geopolitical risks" and gain "Western credibility" suggest that remaining fully China-based imposed an unacceptable cost in terms of access to capital, markets, and technology. This "tax" compels a strategic de-coupling of the product and brand from its country of origin, even if the parent company, Butterfly Effect, maintains a presence in China. This results in a complex, bifurcated operational model.

The relocation to Singapore served as a multi-pronged strategy designed to address these intertwined challenges. Firstly, it facilitates access to US venture capital that would otherwise face severe restrictions under the new regulations, as evidenced by the Benchmark funding and the subsequent US Treasury review. Secondly, Singapore provides better access to international computing resources, including advanced chips, thereby circumventing the impact of US export controls. Thirdly, the move offers refuge from China's "war of a hundred models" [13], where the profit potential in international markets appears more favorable compared to the intensely competitive domestic arena. Fourthly, establishing headquarters in Singapore strategically positions Manus AI as a global company, enhancing its appeal to Western clients and investors and helping it bypass potential sanctions. Finally, the aggressive recruitment in Singapore and the strategic transfer of core Chinese technical staff allow Manus AI to optimize its talent pool by leveraging Chinese engineering expertise while operating from a more globally integrated base.

Singapore's role in this context is not merely as a financial or logistical hub but as an emerging "neutral" ground for AI innovation, specifically for companies caught in geopolitical crosscurrents. The repeated mention of Singapore as a strategic base for Chinese-origin tech companies, including Shein, TikTok, HeyGen, and WIZ.AI, indicates that its value extends beyond geographical proximity to China. Its stable regulatory environment, strong international ties, and access to global talent and infrastructure make it an attractive "safe harbor" where companies can pursue technological development and

market expansion without the direct political scrutiny associated with either the US or China. This positions Singapore as a critical enabler of a more fragmented, multi-polar AI ecosystem.

## 6 Conclusion

Manus AI's strategic retreat from mainland China represents a complex and multi-faceted response to an increasingly challenging operating environment. While geopolitical pressures, specifically US investment restrictions and export controls on advanced AI chips, undeniably played a significant role by limiting access to crucial capital and indispensable technology, internal market dynamics were equally influential. The hyper-competitive "war of a hundred models" in China, characterized by market saturation and aggressive price wars, coupled with Manus AI's declining user numbers and critiques regarding its perceived lack of unique technological breakthroughs, created a compelling internal impetus for seeking international markets.

The relocation to Singapore therefore signifies a strategic maneuver designed to achieve several critical objectives simultaneously: securing access to vital funding, gaining access to essential computing resources, escaping the intense domestic market saturation, and strategically repositioning Manus AI as a global entity. This approach allows the company to leverage its Chinese talent base while shedding the geopolitical liabilities associated with being fully China-based.

Manus AI's case serves as a microcosm of the broader challenges faced by Chinese AI startups navigating a bifurcated global technology landscape. It highlights a growing trend of "de-sinicization" or the adoption of "third path" strategies, where companies proactively seek to establish operational and brand neutrality outside mainland China to access global markets and capital. This trend suggests a potential fragmentation of the global AI ecosystem, with new innovation hubs emerging in geopolitically neutral territories.

The success of Manus AI's pivot will serve as a critical case study for other Chinese technology firms contemplating similar moves. The "third path" strategy, while inherently costly and complex, may become an increasingly vital pathway for Chinese innovation to achieve global scale and secure necessary resources amidst ongoing geopolitical tensions. Future developments will likely see more Chinese-origin companies adopting hybrid operational models, maintaining some research and development or production links to China while strategically relocating core business functions and brand identity to international hubs like Singapore.

## References

- [1] Akin Gump Strauss Hauer & Feld LLP. "Final Regulations Issued by Treasury Restrict US Investment in Chinese Tech Sector". In: *Akin Gump Data Dive* (Nov. 2024). URL: <https://www.akin-gump.com/data-dive>.

Table 2: Drivers Behind Manus AI’s Strategic Pivot: Geopolitical vs. Market/Internal Factors

Category	Geopolitical Factors	Market/Internal Factors
<b>Investment Access</b>	US Outbound Investment Regulations (Jan 2025): Prohibitions/notification for US investments in Chinese AI, creating compliance risks for US VCs. US Treasury Probe of Benchmark Funding: Scrutiny over US\$75M investment due to Chinese ties [9, 1].	
<b>Resource Access</b>	US Chip Export Controls (April 2025): Restrictions on advanced AI chips (Nvidia H100) vital for training AI [15].	
<b>Market Environment</b>	Perception of “Chinese connections as a risk” in US market.	“War of a Hundred Models”: Intense domestic competition with over 130 LLMs, leading to market saturation and price wars. Competition from Chinese Tech Giants: ByteDance (Coze Space), Baidu (AgentBuilder) introducing rival products.
<b>Product Performance</b>		Declining User Numbers: Monthly active users fell from 20M (March) to 10M (May 2025). Critiques on Innovation: Perceived lack of “real technological breakthroughs,” reliance on third-party LLMs.
<b>Operational Strategy</b>		Operational Efficiency Goals: Company’s stated reason for restructuring. Global Talent Strategy: Aggressive recruitment in Singapore, transfer of core technical staff.

- akingump.com/en/insights/blogs/ag-data-dive/final-regulations-issued-by-treasury-restrict-us-investment-in-chinese-tech-sector (visited on 07/15/2025).
- [2] Aibase 基地. “Manus AI Official Website and Social Media Undergo Changes, Chinese Users May Be Affected”. In: *Aibase* (July 2025). URL: <https://www.aibase.com/news/19629> (visited on 07/15/2025).
  - [3] AINVEST. “NVIDIA’s Balancing Act: Navigating China’s AI Market Amid Geopolitical Cross-currents”. In: *AINVEST News* (July 2025). URL: <https://www.ainvest.com/news/nvidia-balancing-act-navigating-china-ai-market-geopolitical-crosscurrents-2507/> (visited on 07/15/2025).
  - [4] Tech in Asia. “Manus shifts HQ to Singapore, cuts China jobs”. In: *Tech in Asia* (July 2025). URL: <https://www.techinasia.com/news/manus-shifts-hq-singapore-cuts-china-jobs> (visited on 07/15/2025).
  - [5] Matthias Bastian. “The startup behind Manus AI shuts down its entire China team to reduce geopolitical risks”. In: *The Decoder* (July 2025). URL: <https://the-decoder.com/the-startup-behind-manus-ai-shuts-down-its-entire-china-team-to-reduce-geopolitical-risks/> (visited on 07/15/2025).
  - [6] Baytech Consulting. “Manus AI: An Analytical Guide to the Autonomous AI Agent 2025”. In: *Baytech Consulting Blog* (May 2025). URL: <https://www.baytechconsulting.com/blog/manus-ai-an-analytical-guide-to-the-autonomous-ai-agent-2025> (visited on 07/15/2025).
  - [7] Bloomberg Quicktake. *Chinese AI Agent Startup Manus Shifts Headquarters To Singapore*. 2025. URL: <https://www.youtube.com/watch?v=dLB0qQwjEs8> (visited on 07/15/2025).
  - [8] Channel NewsAsia. “Chinese firm behind AI agent Manus relocates to Singapore amid US chip curbs”. In: *Channel NewsAsia* (July 2025). URL: <https://www.channelnewsasia.com/east-asia/chinese-firm-manus-ai-relocates-singapore-us-chip-curbs-nvidia-5229391> (visited on 07/15/2025).
  - [9] Fox Rothschild LLP. “Investments in Chinese Technology Companies Limited by New US Outbound Investment Rule”. In: *Fox Rothschild LLP Publications* (Jan. 2025). URL: <https://www.foxrothschild.com/publications/investments-in-chinese-technology-companies-limited-by-new-us-outbound-investment-rule> (visited on 07/15/2025).
  - [10] Insights For Success. “Manus AI: China’s Bold Step Forward—Promises and Challenges”. In: *Kiledjian.com* (Mar. 2025). URL: <https://www.kiledjian.com/main/2025/3/15/manus-ai-chinas-bold-step-forwardpromises-and-challenges> (visited on 07/15/2025).
  - [11] Han Jing. “China’s AI agent Manus relocates HQ to Singapore”. In: *City News Service* (July 2025). URL: <https://www.citynewsservice.cn/news/China's-AI-agent-Manus-relocates-HQ-to-Singapore-7kr585dn> (visited on 07/15/2025).
  - [12] Yee Loon. “Manus AI relocates to Singapore and trims China workforce amid US chip export scrutiny”. In: *The Online Citizen* (July 2025). URL: <https://www.theonlinecitizen.com/2025/07/10/manus-ai-relocates-to-singapore-and-trims-china-workforce-amid-us-chip-export-scrutiny/> (visited on 07/15/2025).
  - [13] Perplexity AI. “China’s War of a Hundred Models”. In: *Perplexity AI* (2025). URL: <https://www.perplexity.ai/page/china-s-war-of-a-hundred-model-DDDsUWBuRDyLkUWBF0p0kQ> (visited on 07/15/2025).
  - [14] STAFF REPORTER. “Chinese developer of Manus AI restructures amid layoff rumors”. In: *The Standard* (July 2025). URL: <https://www.thestandard.com.hk/tech-and-startup/article/306349/Chinese-developer-of-Manus-AI-restructures-amid-layoff-rumors> (visited on 07/15/2025).


- [15] Times of India. “Nvidia may release new AI chip designed specifically for China”. In: *Times of India* (July 2025). URL: <https://timesofindia.indiatimes.com/technology/tech-news/nvidia-may-release-new-ai-chip-designed-specifically-for-china/articleshow/122366728.cms> (visited on 07/15/2025).
- [16] World Economic Forum. “Why China’s AI breakthroughs should come as no surprise”. In: *World Economic Forum* (June 2025). URL: <https://www.weforum.org/stories/2025/06/china-ai-breakthroughs-no-surprise/> (visited on 07/15/2025).



# A Comprehensive Review of Qwen3-Coder: Official Capabilities, Benchmarks, and Community Insights

Neruthes  
Gemini (Google)

2025-07-24

 Article copyright licensing scheme: **Public Domain.**

## Abstract

This literature review provides an in-depth analysis of Qwen3-Coder, the latest large language model from the QwenLM Team, focusing on its official announcement and initial community reception. The analysis synthesizes key architectural innovations, advanced training paradigms—including novel reinforcement learning strategies—and claimed state-of-the-art benchmark performances in agentic coding, browser-use, and tool-use. Concurrently, it critically examines the community’s immediate concerns, particularly revolving around the formidable hardware requirements for local deployment and the efficacy of various quantization techniques. The review highlights the model’s significant advancements in context handling and multi-turn problem-solving, while also addressing practical drawbacks such as resource intensity and ongoing discussions regarding benchmark transparency and real-world reliability. Finally, concrete directions for future improvements are proposed, emphasizing accessibility, robust validation, and ecosystem development to maximize Qwen3-Coder’s impact within the software development landscape.

## 1 Introduction

### 1.1 Background and Significance of Qwen3-Coder

The recent announcement of Qwen3-Coder by the QwenLM Team marks a pivotal moment in the evolution of large language models (LLMs) specifically tailored for code generation and complex agentic tasks.[5, 17] This model is positioned as a significant

leap in open-source artificial intelligence for software development, aiming to redefine automated software engineering.[24, 2] Its introduction into the competitive landscape, alongside proprietary models such as OpenAI’s GPT-4o, Anthropic’s Claude 3.5 Sonnet, and Google’s Gemini 1.5 Pro, as well as open-source rivals like DeepSeek-Coder V2 and Meta’s Code Llama 70B, underscores the rapid advancements and increasing demand for highly capable coding LLMs.[5, 11, 4, 6, 13]

The consistent emphasis in official announcements and various reviews on Qwen3-Coder’s “agentic” capabilities and its direct comparison to top-tier proprietary models indicates a strategic positioning within a maturing market.[17, 6, 30, 1, 16] This release is not merely a technical unveiling; it represents a deliberate attempt to establish a new standard for autonomous software development. By highlighting agentic features and adopting an open-source approach, the QwenLM Team is not simply releasing a new model, but is actively challenging the perception that only closed-source models can achieve such sophistication. The focus on “agentic” capabilities suggests a shift in the model’s value proposition from basic code generation to more advanced code problem-solving and workflow automation, which holds higher utility for developers and enterprises.

## **1.2 Scope and Objectives of the Review**

This review synthesizes information from the official Qwen3-Coder blog post [17] and initial community discussions, primarily sourced from Hacker News [7, 20], to provide a balanced perspective on the model. The objectives of this analysis include: detailing the model’s architectural innovations and unique training methodologies; analyzing its claimed benchmark performance across various coding and agentic tasks; identifying and discussing the practical challenges and drawbacks raised by the community, particularly concerning local deployment and resource intensity; and proposing future directions for model development and ecosystem enhancement.

# **2 Qwen3-Coder: Architectural Innovations and Training Paradigms**

## **2.1 Model Architecture and Key Specifications (MoE, Parameters, Context Length)**

The flagship variant, Qwen3-Coder-480B-A35B-Instruct, is characterized as a Mixture-of-Experts (MoE) model. It features a substantial 480 billion total parameters, with only 35 billion active parameters during inference.[5, 17, 24, 2, 13, 16, 19, 9, 21, 12, 31, 29] This MoE design is a critical innovation, engineered to balance computational efficiency with high performance.[2, 1]

A particularly noteworthy feature is its native support for a 256K token context

length, which can be extended up to 1 million tokens through extrapolation methods such as YaRN.[5, 17, 24, 2, 6, 13, 16, 19, 9, 21, 12, 31, 29] This extensive context window is specifically optimized for handling repo-scale and dynamic data, which is crucial for complex agentic coding tasks.[17, 6, 12]

The combination of a large MoE model and an enormous context window is a deliberate design choice. The MoE architecture enables computational efficiency despite the massive total parameter count by activating only a subset of experts per token.[2, 1] This efficiency is paramount because agentic coding tasks, especially those requiring “repo-scale” understanding, inherently necessitate processing vast amounts of contextual information.[17, 24, 6, 12] A large context window without efficient inference would be prohibitively expensive or slow in practical applications. Consequently, the MoE architecture facilitates the practical application of such an expansive context for complex, multi-file, and multi-turn coding scenarios, directly supporting the model’s ambitious “agentic” capabilities and addressing the computational overhead associated with large context windows.

## **2.2 Pre-training Advancements: Scaling Tokens, Context, and Synthetic Data**

Qwen3-Coder’s development involved pre-training on an impressive 7.5 trillion tokens, with a substantial 70% code ratio. This extensive dataset ensures robust coding capabilities while preserving general and mathematical abilities.[17, 6] This massive training corpus is a key factor in achieving high code quality and the model’s ability to handle diverse coding tasks.[6]

The model also benefits from advancements in context scaling, building upon previous Qwen models that extended context using techniques like YaRN.[6, 14, 15] Furthermore, synthetic data scaling played a significant role. By leveraging Qwen2.5-Coder to refine and rewrite noisy data, the team achieved substantial improvements in overall data quality. This process is analogous to providing tailored “practice exercises” specifically designed to enhance the model’s skills.[6, 29]

The emphasis on a high “70% code ratio” within the 7.5 trillion tokens, coupled with the use of “Synthetic Data Scaling” via Qwen2.5-Coder, indicates a sophisticated understanding that mere data quantity is insufficient. The practice of using a previous model to refine noisy data underscores a strong commitment to data quality and relevance for coding tasks. This distinction is crucial, as high-quality, synthetically generated data can significantly amplify the effectiveness of a vast token count, leading to more robust and capable models, rather than simply scaling up on potentially noisy or less relevant information. This approach directly addresses the challenge of ensuring high-quality input for model training, mitigating potential issues of low-quality output.

## 2.3 Post-training Innovations: Code RL and Long-Horizon Agent RL

Qwen3-Coder incorporates innovative reinforcement learning (RL) strategies in its post-training phase:

- **Scaling Code RL (“Hard to Solve, Easy to Verify”)**: This approach centers on execution-driven, large-scale reinforcement learning applied to a broad set of real-world coding tasks. By automatically scaling test cases for diverse coding challenges, the development team created high-quality training instances, which significantly boosted code execution success rates and yielded benefits for other tasks.[17] This methodology directly contributes to the generation of more reliable and functional code.[6]
- **Scaling Long-Horizon RL (Agent RL)**: This strategy was designed to enable the model to solve complex, multi-turn software engineering tasks, such as those found in SWE-Bench, through continuous interaction with an environment. This involves planning, utilizing tools, receiving feedback, and making iterative decisions.[17, 29] A scalable system, capable of running 20,000 independent environments in parallel on Alibaba Cloud’s infrastructure, facilitated this process. This infrastructure provided the necessary feedback for large-scale reinforcement learning and supported evaluation at scale.[17, 29]

The detailed description of running “20,000 independent environments in parallel on Alibaba Cloud’s infrastructure” for Long-Horizon RL reveals the immense engineering effort and computational resources invested in training truly agentic models. This level of parallelization for environment interaction and feedback represents a significant barrier to entry for many research groups and highlights Alibaba’s substantial commitment to this area. It suggests that achieving advanced “agentic” capabilities is not solely dependent on model architecture or data, but also on sophisticated, large-scale infrastructure and methodologies capable of simulating complex, real-world problem-solving loops. This capability provides a competitive advantage that directly contributes to the model’s claimed state-of-the-art performance in agentic tasks.

## 2.4 Open-Source Tools and Ecosystem Integration

In conjunction with the model release, the QwenLM Team open-sourced “Qwen Code,” a command-line interface (CLI) tool designed for agentic coding. This tool is explicitly noted as being forked and adapted from Gemini Code.[5, 17, 12, 31, 29, 15]

Qwen Code supports the OpenAI SDK for calling LLMs, which signifies a deliberate focus on interoperability and ease of integration into existing developer workflows.[17, 15] The model is also designed to work seamlessly with other popular tools such as Claude Code and Cline.[5, 17, 1, 12, 31, 29]

Open-sourcing Qwen Code and ensuring compatibility with widely used interfaces like the OpenAI SDK, Claude Code, and Cline represents a strategic move to foster broader adoption.[5, 17, 1, 12, 31, 29, 15] By providing familiar tools and interfaces, QwenLM effectively lowers the barrier for developers to experiment with and integrate Qwen3-Coder into their existing development environments. This approach is particularly important for a large model that might otherwise face significant deployment hurdles. It acknowledges that the widespread impact of a model depends not only on its inherent capabilities but also on the robustness and accessibility of the surrounding ecosystem.

## **3 Benchmarking Performance and State-of-the-Art Claims**

### **3.1 Agentic Coding, Browser-Use, and Tool-Use Performance**

The official announcement asserts that Qwen3-Coder-480B-A35B-Instruct “sets new state-of-the-art results among open models on Agentic Coding, Agentic Browser-Use, and Agentic Tool-Use”. [17, 6, 1, 16, 19, 9, 21, 12, 31, 29] Furthermore, it is claimed to achieve performance “comparable to Claude Sonnet 4”. [17, 11, 6, 1, 16, 19, 9, 21, 12, 31, 29]

The repeated assertion of “SOTA among open models” alongside “comparable to Claude Sonnet 4” creates a compelling narrative. [17, 11, 6, 1, 16, 19, 9, 21, 12, 31, 29] This positions Qwen3-Coder not merely as the leading open-source option, but as a credible, potentially more cost-effective alternative to prominent proprietary models. This directly addresses the community’s interest in “Local vs. Cloud Models” [7], suggesting that users may no longer need to accept significant compromises on performance when opting for an open-source solution that can be deployed locally (with appropriate quantization).

### **3.2 SWE-Bench Verified and Other Code-Centric Benchmarks**

On the SWE-Bench Verified benchmark, Qwen3-Coder is reported to achieve state-of-the-art performance among open-source models *without test-time scaling*. [5, 17, 24, 2, 9, 21, 12, 31, 29] This achievement is presented as evidence of its robust long-horizon RL capabilities. [17] The model also demonstrates strong performance on other coding benchmarks, leading on CodeForces ELO, BFCL, and LiveCodeBench v5. [24, 2, 3] For example, a related Qwen3 model (Qwen3-235B) scored 2056 on CodeForces ELO, surpassing DeepSeek-R1 and Gemini 2.5 Pro. [3] Additionally, Qwen3-Coder achieves 61.8% on Aider Polygot [27] and a Terminal-Bench accuracy of 37.5% for the Qwen3-

Coder-480A35 variant.[15]

The phrase “without test-time scaling” on SWE-Bench Verified is a critical detail.[5, 17, 24, 2, 9, 21, 12, 31, 29] SWE-Bench evaluates AI agents on real-world bug-fixing tasks.[8, 28] “Test-time scaling” typically refers to techniques such as multiple inference attempts or complex prompting strategies employed during evaluation to artificially inflate scores. By achieving state-of-the-art performance without such scaling, QwenLM implies a more inherent and robust capability, directly attributable to its Long-Horizon RL training. However, the varying scores across different benchmarks (e.g., 61.8% on Aider Polygot versus 37.5% on Terminal-Bench) suggest that while the model is strong in specific agentic tasks, its performance is not uniformly dominant across all coding challenges. This indicates that “agentic coding” is a multifaceted domain, and leading performance in one sub-area does not necessarily translate to leading performance across the entire spectrum.

### **3.3 Comparative Analysis with Proprietary Models (e.g., Claude Sonnet 4)**

Qwen3-Coder is claimed to be “comparable to Claude Sonnet 4”. [17, 11, 6, 1, 16, 19, 9, 21, 12, 31, 29] Some users have even reported it to be “much faster than Claude Sonnet 4 with similar results”. [7] On the TAU-Bench Retail benchmark, Qwen3-Coder notably outperforms Claude Sonnet 4. [6] However, a comparison involving Qwen3 32B (a different model variant, not the Coder-480B) showed Claude Sonnet 4 outperforming it in AIME 2025 (85.0% vs. 72.9%), while being significantly more expensive for both input and output tokens. [11]

## **4 Community Reception and Practical Deployment Challenges**

### **4.1 The Imperative of Local Deployment and Quantization Efforts**

Initial community reception, particularly on Hacker News, immediately converged on the practical implications of deploying such a large model. [7] This intense focus highlights a strong desire among users to run LLMs locally, driven by concerns over cost, data privacy, and compliance. [24, 11, 7, 20, 27, 3, 18]

### **4.2 Hardware Requirements and Inference Performance**

Running the Qwen3-Coder-480B-A35B-Instruct locally imposes substantial hardware demands. For a dynamic 2-bit quantization, the model requires 24GB of VRAM and

Table 1: Qwen3-Coder’s Claimed Benchmark Performance vs. Key Competitors

Benchmark Model	/ <b>Qwen3-Coder-480B-A35B-Instruct</b>	Claude Sonnet 4	GPT-4.1	Kimi K2	DeepSeek-Coder V2	Gemini 2.5 Pro
Agentic Coding	SOTA (Open Models) [17]	Comparable [17]	-	-	-	-
Agentic Browser-Use	SOTA (Open Models) [17]	Comparable [17]	-	-	-	-
Agentic Tool-Use	SOTA (Open Models) [17]	Comparable [17]	-	-	-	-
SWE-Bench Verified (no test-time scaling)	SOTA (Open Models) [17]	-	-	-	-	-
CodeForces ELO	Lead [2]	-	-	-	-	-
BFCL	Lead [2]	-	-	-	-	-
LiveCodeBench v5	Lead [2]	-	-	-	-	-
Aider Polygot	61.8% [27]	-	-	-	-	-
Terminal-Bench	37.5% [15]	-	-	-	-	-
TAU-Bench Re-tail	Re-Outperforms Claude Sonnet 4 [6]	-	-	-	-	-

*Note: “SOTA” refers to State-of-the-Art among open models. “Comparable” refers to official claims of parity with proprietary models. Numerical scores are provided where available for the specific Qwen3-Coder variant or related Qwen3 models.*

128GB of RAM. For 4-bit quantization, approximately 250GB of RAM is needed, escalating to around 500GB for FP8.[7] Inference speed is significantly influenced by RAM bandwidth, with recommendations for workstations equipped with 8-channel DDR5 memory to optimize performance.[7] Estimated speeds for a setup comprising a 24GB GPU and 128GB RAM are between 3-5 tokens per second, which can drop to less than 1 token/s if RAM capacity is insufficient.[7] Despite these constraints, some users have reported satisfactory performance at approximately 1.5 tokens/second.[7, 18]

The community’s immediate focus on local deployment and quantization, while indicative of a strong desire for self-hosting powerful LLMs, also reveals a significant gap between this “local dream” and the “hardware reality” for most individual developers or even smaller teams. While GPUs with 24GB VRAM (such as the RTX 4090) are considered consumer-grade, the accompanying RAM requirements push into workstation or server-class hardware territory. This implies that while local deployment is technically feasible, it remains largely inaccessible for the average user without substantial investment, creating a practical barrier to widespread adoption despite the model’s open-source nature. This tension between aspiration and practical limitation is a key challenge for broader accessibility.

### 4.3 Dynamic Quantization: Technical Nuances and Community Adoption

Efforts by community members, notably ‘unsloth’, to create quantized versions (e.g., 2-bit to 8-bit GGUFs) for local execution have been a prominent topic of discussion.[7, 27] A primary concern revolved around the viability of highly aggressive quantizations, such as pure 2-bit, with some users reporting previous negative experiences where such low-bit quantizations resulted in “completely broken” models. However, ‘danielhanchen’ from Unsloth provided clarification on their “dynamic quantization” approach, explaining that it involves a mixture of 2, 3, 4, 5, 6, and 8-bit precision. In this method, “important layers are in 8bit, 6bit. Less important ones are left in 2bit”. [7] This intelligent quantization strategy, which involves inspecting activation and weight quantization errors, has been recognized as a crucial advancement in model compression.[7] The process of dynamically quantizing a model of Qwen3-Coder’s scale is itself resource-intensive, requiring several hours and significant cloud computing resources.[7]

The detailed discussion surrounding skepticism about 2-bit quantization and the subsequent explanation of Unsloth’s “dynamic quantization” highlights that model compression is not a simple, universally applied solution; rather, it is an evolving art. The concept of identifying “important layers” and dynamically applying varied precision suggests a sophisticated and ongoing area of research, far from a mature, fully solved problem. The fact that the quantization process itself demands significant resources and time implies that while it enables local inference, the *creation* of these optimized models remains a bottleneck. This often necessitates centralized cloud resources for the initial optimization effort, meaning that accessibility, while improved for inference, is not yet



fully decentralized in terms of model preparation.

## 4.4 Real-World Application and Productivity Debates

Discussions within the community extend to the practical impact of agentic coding on software engineering workflows. Users actively debate whether these tools genuinely enhance productivity, particularly for tasks beyond direct code generation.[7] Agentic coding tools are being applied to automate non-coding overhead tasks, such as writing Git commit messages, creating or updating tickets, and summarizing meetings.[7] Some users have even found AI-written Git messages and tickets to be superior to their manually crafted versions.[7]

A notable debate emerged concerning the reported low percentage of time software engineers spend on “making code changes” (cited as 5% in one user’s breakdown). Some community members characterized this as a “serious organizational dysfunction,” while others contended that it represents a strategic feature for large technology companies, allowing engineers to focus on maintenance and speculative feature development.[7]

The debate regarding the minimal time spent by software engineers on “making code changes” and the application of agentic tools to “non-coding overhead” signifies a fundamental shift in the perception of “developer productivity” in the AI era.[7] If AI can automate these ancillary, yet time-consuming, tasks, the value proposition of a human developer may shift from raw coding output to higher-level activities such as design, architecture, and complex problem-solving, or even the management of AI agents. This suggests that the impact of agentic AI might be less about replacing human coders and more about redefining the coding role and the broader software development lifecycle, potentially freeing human engineers for more complex, creative, or strategic work.

## 4.5 Concerns: Hallucination, Context Handling, and Reliability

Users frequently express concerns regarding LLMs “hallucinating” code or information, particularly for less mainstream or complex tasks.[7, 20, 18] This highlights the ongoing necessity for human oversight and careful prompting to ensure reliable outputs.[7] Some users specifically reported instances where Qwen3-Coder appeared to ignore system prompts, struggled with context, and exhibited rigid tool calls, giving the impression of “formulaic” responses rather than adaptive problem-solving.[18] One user noted hallucination issues specifically when the model was engaged in code-related tasks, despite its satisfactory performance on other types of prompts.[20]

Furthermore, the proliferation of various agentic coding tools and models has led to a perceived “ridiculous” situation of maintaining separate configuration files (e.g., CLAUDE.md, MISTRAL.md, QWEN.md) within repositories. This fragmentation has generated a strong desire within the community for greater standardization of agent con-

figuration protocols.[7]

Table 2: Summary of Community-Identified Drawbacks and Proposed Solutions

Drawback	Cate- gory	Specific Issue	Community Ob- servation / Impact	Proposed / Exist- ing Solutions	Relevant Snippet IDs
<b>Resource Inten- sity</b>		High VRAM/RAM requirements for local inference	Limits accessibil- ity for individual developers and smaller teams; significant initial hardware invest- ment	Dynamic quanti- zation (Unsloth), multi-GPU setups, MoE offloading strate- gies	[7, 27]
<b>Hallucination / Reliability</b>		Inconsistent or in- correct code / in- formation genera- tion	Requires hu- man oversight; reduces trust in autonomous capabilities; can lead to debugging effort	Careful prompt- ing; potential for self-improvement in future models	[7, 20, 18]
<b>Context Han- dling</b>		Struggles with context in multi- file projects; ignores system prompts	Diminishes effectiveness in complex, repo-scale tasks; requires more manual interven- tion	Improved model training for contextual under- standing; better prompt engineer- ing by users	[18]
<b>Tool Call Rigid- ity</b>		“Rigid” or “for- mulaic” tool calls; lacks adaptive “thinking”	Limits flexi- bility in novel scenarios; sug- gests template- filling over true problem-solving	Refined post- training tech- niques; enhanced instruction fol- lowing	[18]
<b>Workflow Frag- mentation</b>		Proliferation of model-specific configuration files	Creates mainte- nance burden; hinders seamless integration of multiple agents	Standardization efforts (e.g., AGENTS.md protocol); sym- linking	[7]

## 5 Critical Assessment: Drawbacks and Areas for Improvement

### 5.1 Resource Intensity and Accessibility Barriers

The most significant drawback of Qwen3-Coder is the sheer size of its 480 billion parameter model, which poses substantial challenges for local deployment and widespread

accessibility without advanced compression techniques.[7] While dynamic quantization represents a promising step towards reducing the model’s footprint, the remaining hardware requirements—such as 24GB VRAM and 128GB RAM even for 2-bit quantized versions—mean that full-precision inference, or even highly quantized inference, remains beyond the reach of many individual developers and smaller teams.[7]

Furthermore, the process of dynamic quantization itself is not trivial; it is resource-intensive, requiring significant cloud computing resources and several hours (estimated at 8 hours minimum for Qwen3-Coder-480B) to complete.[7] This indicates that even the production of these more accessible versions is a complex undertaking. This situation presents a paradox: while the open-source release and quantization efforts aim for decentralized access through local deployment, the sheer scale of the model implies that the *creation* of these accessible quantized versions, and certainly the initial training, remains highly centralized and resource-intensive. This creates a dependency on specialized entities, such as Unsloth or Alibaba Cloud, for the broader community to effectively leverage the model. True democratization of such large models necessitates not only open weights but also democratized means of optimization and deployment.

## 5.2 Benchmark Transparency and Reproducibility Concerns

As with any newly released model, the establishment of detailed and independently verifiable benchmarks across a wider, more diverse range of real-world coding and agentic tasks would significantly strengthen Qwen3-Coder’s standing. Concerns regarding “deceptive benchmark hacking” and skepticism about state-of-the-art claims have been voiced within the community, with some users advising against relying solely on benchmarks released by model-developing companies.[30, 22]

Specifically, authors of the Arc AGI benchmark reportedly could not reproduce Qwen’s claimed 41% score.[30, 22] While QwenLM denies engaging in benchmark manipulation [22], community observations indicate that Qwen recently modified their “cradle” (evaluation environment) and enabled tool use, which resulted in a significant (30%) jump in scores for all models, including smaller open ones.[23] These observations raise important questions about the methodology and comparability of benchmarks across different models and evaluation setups. This situation highlights an “arms race” in LLM benchmarking, where companies may optimize their models and evaluation environments to perform exceptionally well on specific, publicly known benchmarks, potentially at the expense of generalizability or real-world robustness. This practice can erode community trust and complicate objective model comparisons. The lack of independent verification and the ease with which benchmark scores can be influenced by subtle methodological changes (e.g., “cradle” adjustments, tool use enablement) suggest that raw benchmark numbers alone are insufficient indicators of a model’s true capabilities. This underscores the need for more standardized, transparent, and independently verifiable evaluation protocols within the broader LLM community.

### 5.3 Model Consistency and Adaptability in Diverse Scenarios

Despite the claimed agentic capabilities, some users have reported issues with Qwen3-Coder, including instances where it appeared to ignore system prompts, struggled with context in multi-file projects, and produced “rigid” or “formulaic” tool calls.[18] These observations suggest a potential lack of adaptive “thinking” beyond simple template filling. One user also experienced hallucination issues when working on code, even while the model performed well on other tasks.[20]

These reported inconsistencies indicate that while the model may excel in specific, structured benchmark scenarios, its real-world performance for complex, nuanced, or novel coding tasks might still necessitate significant human intervention and careful prompting. The reported issues of “ignoring system prompts,” “struggling with context,” and “rigid tool calls” appear to contradict the official narrative of Qwen3-Coder being the “most agentic” model.[18] While the model might perform well on structured agentic benchmarks, these user experiences suggest a gap between *agentic behavior* (executing multi-step tasks with tools) and true *autonomy* or *robust adaptability*. A truly autonomous agent would not disregard instructions or struggle with dynamic context. This implies that the model, despite its advanced reinforcement learning training, may still exhibit brittleness when confronted with real-world complexities that deviate from its training distribution.

## 6 Future Directions and Recommendations

### 6.1 Advancements in Model Compression and Efficient Inference

Continued research into more efficient and less lossy compression techniques, building upon the foundation of dynamic quantization, is crucial. This includes exploring novel quantization methods, sparse model architectures, and highly optimized inference frameworks.[7, 27] A key focus should be on optimizing these models for more common consumer-grade hardware configurations, such as GPUs with 16GB VRAM paired with more modest RAM capacities. This would significantly broaden accessibility for individual developers and smaller teams who lack extensive cloud resources.[7, 27, 18] Further development of Mixture-of-Experts (MoE) offloading strategies to efficiently distribute the computational load across heterogeneous hardware (CPU/GPU) is also recommended.[27]

The persistent emphasis on optimizing for local deployment and the community’s desire for smaller, more optimized variants highlight a significant “democratization bottleneck” for large models, even open-source ones.[7, 27, 18] The widespread success and impact of such powerful models depend not only on their inherent capabilities but

also on their *accessibility*. Overcoming this bottleneck requires continuous innovation in compression and efficient inference, making it feasible for a broader user base to run and experiment with these models without prohibitive hardware investments. This is a critical factor for fostering a vibrant and inclusive open-source ecosystem.

## 6.2 Enhancing Benchmark Rigor and Independent Validation

To cultivate greater community trust and facilitate clearer comparisons, future model releases should prioritize transparent and independently verifiable benchmarks across a wider, more diverse range of real-world coding and agentic tasks.[30, 22] This entails publishing detailed methodologies, comprehensive training data, and evaluation scripts to enable full reproducibility by third parties.[22] Collaboration with independent benchmarking organizations, such as those responsible for SWE-Bench Verified and Terminal-Bench, to conduct and publish results would significantly enhance the credibility of performance claims.[8, 28, 25, 10, 26]

The skepticism surrounding company-released benchmarks and the specific issues related to Arc AGI reproducibility point to a trust deficit within the LLM community.[30, 22] To address this, model developers need to move beyond simply publishing scores. Full transparency in methodology, data, and evaluation scripts, coupled with active engagement with independent evaluators, is essential. This approach would shift the focus from merely “claiming state-of-the-art” to “demonstrating robust, verifiable performance,” which is crucial for long-term adoption and maintaining scientific integrity.

## 6.3 Fostering Community Tooling and Smaller, Optimized Variants

Continued development and support for tools like Qwen Code, coupled with active encouragement of community contributions, will enhance the model’s usability and integration into diverse development environments.[5, 17, 12, 31, 29, 15] Addressing the issue of “configuration file proliferation” through standardized agent configuration protocols (e.g., an ‘AGENTS.md’ standard, as suggested by the community) would significantly improve the developer experience.[7] While the 480B model is undoubtedly powerful, the development of smaller, highly optimized variants that retain a significant portion of its agentic capabilities could broaden its applicability and reduce computational overhead for specific use cases.[27, 18] This strategy would cater to the “good enough” philosophy for certain tasks, where optimal performance is less critical than accessibility and efficiency.[18]

The community’s frustration with tool fragmentation and the desire for smaller models underscore that a model’s true value extends beyond its raw performance; it also encompasses its *integrability* and *usability* within a broader ecosystem.[7, 27, 18] By ac-

tively fostering community tooling, supporting standardization efforts, and developing a range of model sizes, QwenLM can significantly amplify the impact of Qwen3-Coder. This strategy acknowledges that a powerful model alone is insufficient; it must be embedded within a supportive, user-friendly environment to achieve widespread adoption and utility.

## 6.4 Improving Robustness and Reliability for Agentic Workflows

Addressing reported issues of hallucination, the model ignoring system prompts, and rigid tool calls is paramount for widespread real-world agentic adoption.[20, 18] This may necessitate refining post-training techniques, enhancing instruction following capabilities, and improving contextual understanding for multi-file and long-horizon tasks. Further research into self-improvement mechanisms for coding agents, as hinted by QwenLM, could lead to models that autonomously learn from their failures and adapt to novel scenarios.[17] This would bridge the existing gap between merely “agentic” behavior and true “autonomous” intelligence.

The practical struggles reported by the community regarding hallucination and rigid behavior suggest that even models achieving state-of-the-art agentic benchmarks still exhibit a qualitative difference between performing well on structured tests and reliably navigating the messy, unpredictable nature of real-world software engineering.[20, 18] The pursuit of “self-improvement” is critical here.[17] It signifies a shift from training models for specific, predefined tasks to training them for continuous learning and adaptation, which is a hallmark of true intelligence and autonomy. This long-term vision is necessary to overcome current limitations and fully deliver on the promise of truly transformative AI coding assistants.

## 7 Conclusion

Qwen3-Coder represents a significant leap in open-source code-centric large language models, particularly in its agentic capabilities and extensive context handling. Its innovative training methodologies, especially in large-scale Code RL and Long-Horizon Agent RL, position it as a strong contender in the competitive landscape of AI for software development. The model’s claimed state-of-the-art performance on benchmarks like SWE-Bench Verified, alongside its comparability to proprietary models such as Claude Sonnet 4, underscores its technical prowess.

However, the immediate community engagement highlights substantial practical challenges associated with its deployment. These challenges primarily stem from its formidable size and the resulting hardware requirements for local inference. While dynamic quantization offers a promising avenue for accessibility, it also reveals the ongoing research and resource intensity required for effective model compression. Furthermore,

concerns regarding benchmark transparency and the model's consistency in complex, real-world scenarios necessitate continued efforts in independent validation and robustness improvements.

The tension between the model's advanced capabilities and its practical accessibility underscores a critical theme in the current LLM landscape: the democratization of powerful AI. Future advancements will depend not only on pushing the boundaries of model intelligence but also on developing more efficient deployment strategies, fostering a robust open-source ecosystem, and ensuring rigorous, transparent evaluation. Qwen3-Coder, with its blend of scale, innovation, and open-source commitment, sets a new standard, but its ultimate impact will hinge on how effectively these challenges are addressed to empower the broader developer community.

## References

- [1] AInvest. *Alibaba's Qwen3-Coder Model Sets New Standards in Code Modeling with Advanced Agentic Capabilities*. July 23, 2025. URL: <https://www.ainvest.com/news/alibaba-qwen3-coder-model-sets-standards-code-modeling-advanced-agentic-capabilities-2507/> (visited on 07/23/2025).
- [2] Apidog. *Qwen3-Coder is Finally Here and It's Breaking All the Coding Benchmarks*. July 23, 2025. URL: <https://apidog.com/blog/qwen3-coder/> (visited on 07/23/2025).
- [3] DataCamp. *Qwen 3: Features, DeepSeek-R1 Comparison, Access, and More*. July 23, 2025. URL: <https://www.datacamp.com/blog/qwen3> (visited on 07/23/2025).
- [4] Entelligence Blog. *Claude 4 vs Deepseek R1 vs Qwen 3*. July 23, 2025. URL: <https://www.entelligence.ai/blogs/Claude-4-vs-Deepseek-r1-vs-qwen-3> (visited on 07/23/2025).
- [5] Fortune India. *Qwen3-Coder: Alibaba claims world's most advanced agentic AI model for coding*. July 23, 2025. URL: <https://www.fortuneindia.com/technology/qwen3-coder-alibaba-claims-worlds-most-advanced-agentic-ai-model-for-coding/125146> (visited on 07/23/2025).
- [6] Mehul Gupta. *Qwen3-Coder : The best Agentic Code AI, beats Kimi-K2*. July 2025. URL: <https://medium.com/data-science-in-your-pocket/qwen3-coder-the-best-agentic-code-ai-beats-kimi-k2-1f8e6472c42b> (visited on 07/23/2025).
- [7] Hacker News. *Qwen3-Coder: Agentic coding in the world*. July 23, 2025. URL: <https://news.ycombinator.com/item?id=44653072> (visited on 07/23/2025).
- [8] Holistic Agent Leaderboard. *SWE-bench Verified*. July 23, 2025. URL: <https://hal.cs.princeton.edu/swebench> (visited on 07/23/2025).
- [9] Investing.com. *Alibaba unveils Qwen3-Coder-480B, its most powerful coding model*. July 23, 2025. URL: <https://www.investing.com/news/stock-market-news/alibaba-unveils-qwen3coder480b-its-most-powerful-coding-model-93CH-4147019> (visited on 07/23/2025).
- [10] laude-institute. *laude-institute/terminal-bench: A benchmark for LLMs on complicated tasks in the terminal*. July 23, 2025. URL: <https://github.com/laude-institute/terminal-bench> (visited on 07/23/2025).
- [11] LLM Stats. *Claude Sonnet 4 vs Qwen3 32B*. July 23, 2025. URL: <https://llm-stats.com/models/compare/claude-sonnet-4-20250514-vs-qwen3-32b> (visited on 07/23/2025).

- [12] MarkTechPost. *Qwen Releases Qwen3-Coder-480B-A35B-Instruct: Its Most Powerful Open Agentic Code Model Yet*. July 22, 2025. URL: <https://www.marktechpost.com/2025/07/22/qwen-releases-qwen3-coder-480b-a35b-instruct-its-most-powerful-open-agentic-code-model-yet/> (visited on 07/23/2025).
- [13] OpenRouter. *Qwen: Qwen3 Coder –Uptime and Availability*. July 23, 2025. URL: <https://openrouter.ai/qwen/qwen3-coder/uptime> (visited on 07/23/2025).
- [14] Qwen Team. *Qwen2.5-1M Technical Report*. July 23, 2025. URL: [https://qianwen-res.oss-cn-beijing.aliyuncs.com/Qwen2.5-1M/Qwen2\\_5\\_1M\\_Technical\\_Report.pdf](https://qianwen-res.oss-cn-beijing.aliyuncs.com/Qwen2.5-1M/Qwen2_5_1M_Technical_Report.pdf) (visited on 07/23/2025).
- [15] QwenLM. *QwenLM/qwen-code: qwen-code is a coding agent that lives in digital world*. July 23, 2025. URL: <https://github.com/QwenLM/qwen-code> (visited on 07/23/2025).
- [16] QwenLM. *QwenLM/Qwen3-Coder: Qwen3-Coder is the code version of Qwen3, the large language model series developed by Qwen team, Alibaba Cloud*. July 23, 2025. URL: <https://github.com/QwenLM/Qwen3-Coder> (visited on 07/23/2025).
- [17] QwenLM Team. *Qwen3-Coder: Agentic Coding in the World*. July 23, 2025. URL: <https://qwenlm.github.io/blog/qwen3-coder/> (visited on 07/23/2025).
- [18] r/LocalLLaMA. *Kimi K2 vs Qwen3 Coder 480B*. July 23, 2025. URL: [https://www.reddit.com/r/LocalLLaMA/comments/1m6zz1v/kimi\\_k2\\_vs\\_qwen3\\_coder\\_480b/](https://www.reddit.com/r/LocalLLaMA/comments/1m6zz1v/kimi_k2_vs_qwen3_coder_480b/) (visited on 07/23/2025).
- [19] r/LocalLLaMA. *Qwen/Qwen3-Coder-480B-A35B-Instruct*. July 23, 2025. URL: <https://www.reddit.com/r/LocalLLaMA/comments/1m6qc8c/qwenqwen3coder480ba35binstruct/> (visited on 07/23/2025).
- [20] r/LocalLLaMA. *Qwen3- Coder*. July 23, 2025. URL: [https://www.reddit.com/r/LocalLLaMA/comments/1m6mew9/qwen3\\_coder/](https://www.reddit.com/r/LocalLLaMA/comments/1m6mew9/qwen3_coder/) (visited on 07/23/2025).
- [21] r/LocalLLaMA. *Qwen3-Coder is here!* July 23, 2025. URL: [https://www.reddit.com/r/LocalLLaMA/comments/1m6qdet/qwen3coder\\_is\\_here/](https://www.reddit.com/r/LocalLLaMA/comments/1m6qdet/qwen3coder_is_here/) (visited on 07/23/2025).
- [22] r/LocalLLaMA. *Recent Qwen Benchmark Scores are Questionable*. July 23, 2025. URL: [https://www.reddit.com/r/LocalLLaMA/comments/1m6wb5o/recent\\_qwen\\_benchmark\\_scores\\_are\\_questionable/](https://www.reddit.com/r/LocalLLaMA/comments/1m6wb5o/recent_qwen_benchmark_scores_are_questionable/) (visited on 07/23/2025).
- [23] r/singularity. *Kimi K2 is already irrelevant, and it's only been like 1 week. Qwen has updated Qwen-3-235B, and it outperforms K2 at less than 1/4th the size*. July 23, 2025. URL: [https://www.reddit.com/r/singularity/comments/1m5tupt/kimi\\_k2\\_is\\_already\\_irrelevant\\_and\\_its\\_only\\_been/](https://www.reddit.com/r/singularity/comments/1m5tupt/kimi_k2_is_already_irrelevant_and_its_only_been/) (visited on 07/23/2025).
- [24] Gary Svenson. *Qwen3-Coder: The New Titan of AI Coding Models Is Here*. July 2025. URL: <https://garysvenson09.medium.com/qwen3-coder-the-new-titan-of-ai-coding-models-is-here-7cfe7bfc1e> (visited on 07/23/2025).
- [25] tbench.ai. *Terminal-Bench*. July 23, 2025. URL: <https://www.tbench.ai/> (visited on 07/23/2025).
- [26] Terminal-Bench. *Introduction*. July 23, 2025. URL: <https://www.tbench.ai/docs> (visited on 07/23/2025).
- [27] Unsloth Documentation. *Qwen3-Coder: How to Run Locally*. July 23, 2025. URL: <https://docs.unsloth.ai/basics/qwen3-coder-how-to-run-locally> (visited on 07/23/2025).
- [28] Warp. *Warp scores 71% on SWE-bench Verified*. July 23, 2025. URL: <https://www.warp.dev/blog/swe-bench-verified> (visited on 07/23/2025).
- [29] Simon Willison. *Qwen3-Coder: Agentic Coding in the World*. July 22, 2025. URL: <https://simonwillison.net/2025/Jul/22/qwen3-coder/> (visited on 07/23/2025).



- [30] YouTube. *Qwen 3 Coder (480B Tested) + Free APIs + Qwen CLI, Cline, Roo: It's a Good Model but It's Kinda Weird*. July 23, 2025. URL: <https://www.youtube.com/watch?v=kztx5Vkc2-I> (visited on 07/23/2025).
- [31] YouTube. *Qwen Releases Qwen3-Coder-480B-A35B-Instruct: A Leading Open Agentic Code Model*. July 23, 2025. URL: <https://www.youtube.com/watch?v=BQFFcE6B1GM> (visited on 07/23/2025).